

```
In [1]: # Importing necessary libraries and modules
```

```
import numpy as np # NumPy for numerical operations and handling arrays
import matplotlib.pyplot as plt # Matplotlib for plotting and visualization
from keras.models import Sequential # Sequential model for building neural networks
from keras.layers import Dense # Dense layer for fully connected neural networks
from keras.optimizers import Adam # Adam optimizer for training the model
from keras.utils import to_categorical # Utility for one-hot encoding class labels
from keras.layers import Dropout, Flatten # Dropout layer for regularization and Flatten layer
from keras.layers import Conv2D, MaxPooling2D # Convolutional and pooling layers
from keras.models import load_model # Function to load a pre-trained model
import cv2 # OpenCV for image processing tasks
from sklearn.model_selection import train_test_split # Utility to split the dataset
import pickle # Module for serializing and deserializing Python objects
import os # Module for interacting with the operating system (e.g., file operations)
import pandas as pd # Pandas for data manipulation and analysis with DataFrames
from tensorflow.keras.preprocessing.image import ImageDataGenerator # Class for image data preprocessing
```

```
In [2]: # Define the path to the dataset and label file
```

```
path = "Dataset" # Directory containing the image classes
labelFile = 'labels.csv' # CSV file containing labels for the images

# Set parameters for training
batch_size_val = 32 # Number of samples per gradient update
epochs_val = 10 # Number of epochs to train the model
imageDimensions = (32, 32, 3) # Dimensions of the input images (height, width, channels)
testRatio = 0.2 # Ratio of the dataset to be used for testing
validationRatio = 0.2 # Ratio of the training dataset to be used for validation

# Initialize counters and lists to store images and their corresponding class numbers
count = 0 # Class counter
images = [] # List to store images
classNo = [] # List to store class numbers corresponding to the images

# List all subdirectories (classes) in the dataset directory
myList = os.listdir(path) # Get a list of all classes (subdirectories) in the dataset directory
print("Total Classes Detected:", len(myList)) # Print the number of classes detected
noOfClasses = len(myList) # Store the number of classes

print("Importing Classes.....") # Indicate the start of class importing
# Loop through each class directory
for x in range(0, len(myList)):
    myPicList = os.listdir(path + "/" + myList[x]) # List all images in the current class directory
    # Loop through each image in the current class directory
    for y in myPicList:
        curImg = cv2.imread(path + "/" + myList[x] + "/" + y) # Read the image
        images.append(curImg) # Append the image to the images list
        classNo.append(x) # Append the class number to the classNo list
    print(count, end=" ") # Print the current class number
    count += 1 # Increment the class counter

print("\n") # Print a newline for better readability

# Convert the lists of images and class numbers to NumPy arrays for further processing
images = np.array(images) # Convert images list to a NumPy array
classNo = np.array(classNo) # Convert classNo list to a NumPy array
```

Total Classes Detected: 43

Importing Classes.....

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28
31 32 33 34 35 36 37 38 39 40 41 42

```
In [3]: # Split the dataset into training and testing sets
# The test set will be a portion of the images and their corresponding class labels
X_train, X_test, y_train, y_test = train_test_split(images, classNo, test_size=0.2)
# Further split the training set into training and validation sets
# The validation set will be a portion of the training data used to tune the model
X_train, X_validation, y_train, y_validation = train_test_split(X_train, y_train, test_size=0.2)

# Print the shapes of the resulting datasets for verification
print("Data Shapes") # Indicate that data shapes will be printed
print("Train", end=" ") # Print 'Train' without a newline
print(X_train.shape, y_train.shape) # Print the shape of training images
print("Validation", end=" ") # Print 'Validation' without a newline
print(X_validation.shape, y_validation.shape) # Print the shape of validation images
print("Test", end=" ") # Print 'Test' without a newline
print(X_test.shape, y_test.shape) # Print the shape of test images and labels
```

Data Shapes

Train(22271, 32, 32, 3) (22271,)

Validation(5568, 32, 32, 3) (5568,)

Test(6960, 32, 32, 3) (6960,)

```
In [4]: # Load the labels from the CSV file into a DataFrame
data = pd.read_csv(labelFile) # Read the labels from the specified CSV file
print("data shape ", data.shape, type(data)) # Print the shape and type of data

# Initialize a list to store the number of samples for each class
num_of_samples = [] # List to hold the number of samples per class
cols = 5 # Number of columns for displaying images (not used in this snippet)
num_classes = noOfClasses # Total number of classes detected

# Function to convert an image to grayscale
def grayscale(img):
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) # Convert the image from BGR to grayscale
    return img # Return the grayscale image

# Function to equalize the histogram of an image
def equalize(img):
    img = cv2.equalizeHist(img) # Apply histogram equalization to improve contrast
    return img # Return the equalized image

# Function to preprocess an image
def preprocessing(img):
    img = grayscale(img) # Convert the image to grayscale
    img = equalize(img) # Equalize the histogram of the grayscale image
    img = img / 255 # Normalize the pixel values to the range [0, 1]
    return img # Return the preprocessed image

data shape (43, 2) <class 'pandas.core.frame.DataFrame'>
```

```

In [ ]: # Apply the preprocessing function to each image in the training, validation, and test sets
# The preprocessing function converts images to grayscale, equalizes them, and resizes them to 28x28 pixels
X_train = np.array(list(map(preprocessing, X_train))) # Preprocess all training images
X_validation = np.array(list(map(preprocessing, X_validation))) # Preprocess all validation images
X_test = np.array(list(map(preprocessing, X_test))) # Preprocess all test images

# Reshape the training, validation, and test sets to add a channel dimension
# The new shape will be (number of samples, height, width, channels)

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], X_train.shape[2], 1)
X_validation = X_validation.reshape(X_validation.shape[0], X_validation.shape[1], X_validation.shape[2], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], X_test.shape[2], 1)

In [6]: # Create an instance of the ImageDataGenerator class for data augmentation
dataGen = ImageDataGenerator(
    width_shift_range=0.1, # Randomly shift images horizontally by 10%
    height_shift_range=0.1, # Randomly shift images vertically by 10%
    zoom_range=0.2, # Randomly zoom in on images by up to 20%
    shear_range=0.1, # Randomly apply shear transformations by up to 10%
    rotation_range=10, # Randomly rotate images by up to 10 degrees
)

# Fit the data generator to the training data
dataGen.fit(X_train) # This computes the necessary statistics (like mean and standard deviation)

# Create batches of augmented data from the training set
batches = dataGen.flow(X_train, y_train, batch_size=20) # Generate batches of augmented data

# Get the next batch of images and their corresponding labels
X_batch, y_batch = next(batches) # Retrieve one batch of augmented data

In [7]: # Convert the class labels for the training, validation, and test sets to one-hot encoding
y_train = to_categorical(y_train, noOfClasses) # Convert training labels to one-hot encoding
y_validation = to_categorical(y_validation, noOfClasses) # Convert validation labels to one-hot encoding
y_test = to_categorical(y_test, noOfClasses) # Convert test labels to one-hot encoding

```

```
In [ ]: def myModel():
# Initialize a sequential model
model = Sequential()
# Add the first convolutional layer with 60 filters, each of size 5x5
# The input shape is defined as (height, width, channels), where channels is 3
model.add(Conv2D(60, (5, 5), input_shape=(imageDimensions[0], imageDimensions[1], 3)))
# Add a second convolutional layer with 60 filters of size 5x5
model.add(Conv2D(60, (5, 5), activation='relu'))
# Add a max pooling layer to downsample the feature maps
model.add(MaxPooling2D(pool_size=(2, 2)))
# Add a third convolutional layer with 30 filters of size 3x3
model.add(Conv2D(30, (3, 3), activation='relu'))
# Add a fourth convolutional layer with 30 filters of size 3x3
model.add(Conv2D(30, (3, 3), activation='relu'))
# Add another max pooling layer to further downsample the feature maps
model.add(MaxPooling2D(pool_size=(2, 2)))
# Add a dropout layer to reduce overfitting by randomly setting 50% of the input to 0
model.add(Dropout(0.5))
# Flatten the output from the previous layer to feed into the dense layer
model.add(Flatten())
# Add a fully connected (dense) layer with 500 units and ReLU activation
model.add(Dense(500, activation='relu'))
# Add another dropout layer to reduce overfitting
model.add(Dropout(0.5))
# Add the output layer with a number of units equal to the number of classes
model.add(Dense(numOfClasses, activation='softmax'))
# Compile the model with Adam optimizer, categorical_crossentropy loss
model.compile(Adam(learning_rate=0.001), loss='categorical_crossentropy')
# Return the constructed model
return model
```

```
In [ ]: # Create an instance of the model defined in the myModel function
model = myModel()
# Print the summary of the model architecture, including layer types, output shapes, and number of parameters
print(model.summary())
# Train the model using the fit method
# The training data is generated using the data generator with data augmentation
history = model.fit(
    dataGen.flow(X_train, y_train, batch_size=32), # Generate batches of
    steps_per_epoch=len(X_train) // 32, # Number of steps to complete one epoch
    epochs=epochs_val, # Number of epochs to train the model
    validation_data=(X_validation, y_validation), # Validation data to evaluate the model
    shuffle=1 # Shuffle the training data before each epoch
)
```

c:\Users\BABAR\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not use an `input_shape`/`input_dim` argument to a layer. When using Sequential model, prefer using an `Input(shape)` object as the first layer in the model instead of passing `input_shape` to the first layer.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param
conv2d (Conv2D)	(None, 28, 28, 60)	1,560
conv2d_1 (Conv2D)	(None, 24, 24, 60)	90,000
max_pooling2d (MaxPooling2D)	(None, 12, 12, 60)	
conv2d_2 (Conv2D)	(None, 10, 10, 30)	16,200
conv2d_3 (Conv2D)	(None, 8, 8, 30)	8,100
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 30)	
dropout (Dropout)	(None, 4, 4, 30)	
flatten (Flatten)	(None, 480)	
dense (Dense)	(None, 500)	240,500
dropout_1 (Dropout)	(None, 500)	
dense_1 (Dense)	(None, 43)	21,545

Total params: 378,023 (1.44 MB)

Trainable params: 378,023 (1.44 MB)


Non-trainable params: 0 (0.00 B)

None

Epoch 1/10

c:\Users\BABAR\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

695/695  114s 154ms/step - accuracy: 0.1723 - loss: 3.057
val_accuracy: 0.7656 - val_loss: 0.7548

Epoch 2/10

1/695  1:41 147ms/step - accuracy: 0.5938 - loss: 1.728

c:\Users\BABAR\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\trainers\epoch_iterator.py:107: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.

self._interrupted_warning()

695/695 ————— 6s 8ms/step - accuracy: 0.5938 - loss: 1.7286 -
 val_accuracy: 0.7696 - val_loss: 0.7635
 Epoch 3/10

695/695 ————— 103s 148ms/step - accuracy: 0.6018 - loss: 1.290
 val_accuracy: 0.9318 - val_loss: 0.2929
 Epoch 4/10

695/695 ————— 6s 8ms/step - accuracy: 0.6875 - loss: 1.1744 -
 val_accuracy: 0.9307 - val_loss: 0.3226
 Epoch 5/10

695/695 ————— 114s 163ms/step - accuracy: 0.7474 - loss: 0.817
 val_accuracy: 0.9574 - val_loss: 0.1636
 Epoch 6/10

695/695 ————— 6s 9ms/step - accuracy: 0.8438 - loss: 0.6667 -
 val_accuracy: 0.9578 - val_loss: 0.1575
 Epoch 7/10



695/695 ————— 99s 143ms/step - accuracy: 0.8059 - loss: 0.6343
 val_accuracy: 0.9567 - val_loss: 0.1444
 Epoch 8/10

695/695 ————— 5s 8ms/step - accuracy: 0.7500 - loss: 0.7462 -
 val_accuracy: 0.9626 - val_loss: 0.1356
 Epoch 9/10

695/695 ————— 91s 130ms/step - accuracy: 0.8399 - loss: 0.5115
 val_accuracy: 0.9616 - val_loss: 0.1325
 Epoch 10/10

695/695 ————— 5s 7ms/step - accuracy: 0.7812 - loss: 1.0221 -
 val_accuracy: 0.9664 - val_loss: 0.1193

```
In [ ]: # Create a new figure for plotting the loss
plt.figure(1)
# Plot the training loss over epochs
plt.plot(history.history['loss'])
# Plot the validation loss over epochs
plt.plot(history.history['val_loss'])
# Add a legend to differentiate between training and validation loss
plt.legend(['training', 'validation'])
# Set the title of the plot
plt.title('Loss')
# Label the x-axis as 'epoch'
plt.xlabel('epoch')
# Create a new figure for plotting the accuracy
plt.figure(2)
# Plot the training accuracy over epochs
plt.plot(history.history['accuracy'])
# Plot the validation accuracy over epochs
plt.plot(history.history['val_accuracy'])
# Add a legend to differentiate between training and validation accuracy
plt.legend(['training', 'validation'])
# Set the title of the plot
plt.title('Accuracy')
# Label the x-axis as 'epoch'
plt.xlabel('epoch')
# Display the plots
plt.show()
```



```
In [ ]: # Evaluate the model on the test dataset
# The evaluate method returns the loss value and metrics specified during
score = model.evaluate(X_test, y_test, verbose=0)
# Print the test score (loss) from the evaluation
print('Test Score:', score[0])
# Print the test accuracy from the evaluation
print('Test Accuracy:', score[1])
```

Test Score: 0.12101157754659653
Test Accuracy: 0.9643678069114685

```
In [ ]: # Save the model
model.save('model.keras')
model.save('model.h5')
# Load the model
loaded_model = load_model('model.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to built. `model.compile_metrics` will be empty until you train or evaluate the

```
In [ ]: def predict_image(image_path,model,image_dimensions=(32, 32)):
# Load the image
img = cv2.imread(image_path)
if img is None:
    raise ValueError(f"Image at path '{image_path}' not found.")
# Preprocess the image
img = cv2.resize(img, (image_dimensions[0], image_dimensions[1]))
img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) # Convert to grayscale
img = cv2.equalizeHist(img) # Equalize
img = img / 255.0 # Normalize pixel values
img = img.reshape(1, image_dimensions[0], image_dimensions[1], 1) # ,
# Predict using the model
prediction = model.predict(img)
predicted_class = np.argmax(prediction) # Get the class with the high
df =pd.read_csv('labels.csv')
Sign = df.iloc[predicted_class,1]
result = f"The Traffic-Sign in the Image States That: \n{Sign}"
print (result)
```

```
In [ ]: #Predicting a new image
path = "C:/Users/BABAR/Desktop/Traffic Signs Recognition/images.png"
predict_image(path,loaded_model)
```

1/1 ————— 0s 239ms/step
The Traffic-Sign in the Image States That:
Stop

```
In [ ]:
```

```
In [ ]:
```